

# Massively Parallel Tensor Contraction for High-Accuracy Quantum Chemical Calculations

**Andreas Irmeler**

Grüneis group

Vienna University of Technology, Austria



TECHNISCHE  
UNIVERSITÄT  
WIEN

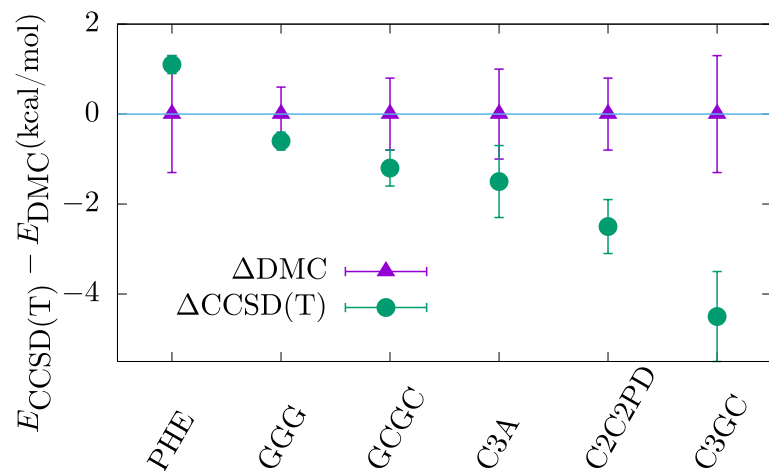
# Accurate ab-initio calculations:

Interactions between large molecules pose a puzzle for reference quantum mechanical methods

Yasmine S. Al-Hamdani, Péter R. Nagy, Andrea Zen, Dennis Barton, Mihály Kállay

Jan Gerit Brandenburg  & Alexandre Tkatchenko

NATURE COMMUNICATIONS (2021)



# Accurate ab-initio calculations:

Interactions between large molecules pose a puzzle  
for reference quantum mechanical methods

Yasmine S. Al-Hamdani<sup>1,2,9</sup>, Péter R. Nagy<sup>3,9</sup>, Andrea Zen<sup>4,5,6,7</sup>, Dennis Barton<sup>2</sup>, Mihály Kállay<sup>3</sup>,  
Jan Gerit Brandenburg<sup>8,10</sup> & Alexandre Tkatchenko<sup>2,10</sup>

NATURE COMMUNICATIONS | (2021)12:3927 | <https://doi.org/10.1038/s41467-021-24119-3> | [www.nature.com/naturecommunications](http://www.nature.com/naturecommunications)

In obtaining CCSD(T) interaction energies, the sources of error  
are:

- Single-particle basis representation of the CCSD(T) wavefunction.
- Local approximations of long-range electron correlation according to the LNO scheme.
- Neglected core electron correlation.
- Missing high-order many-electron contributions beyond CCSD(T).

# Accurate ab-initio calculations:

Interactions between large molecules pose a puzzle for reference quantum mechanical methods

Yasmine S. Al-Hamdani<sup>1,2,9</sup>, Péter R. Nagy<sup>3,9</sup>, Andrea Zen<sup>4,5,6,7</sup>, Dennis Barton<sup>2</sup>, Mihály Kállay<sup>3</sup>, Jan Gerit Brandenburg<sup>8,10</sup> & Alexandre Tkatchenko<sup>2,10</sup>

NATURE COMMUNICATIONS | (2021)12:3927 | <https://doi.org/10.1038/s41467-021-24119-3> | [www.nature.com/naturecommunications](http://www.nature.com/naturecommunications)

## Understanding Discrepancies of Wavefunction Theories for Large Molecules

Tobias Schäfer, Andreas Irmler, Alejandro Gallo, and Andreas Grüneis

arXiv: 2407.01442, Nat. Comm (accepted)

In obtaining CCSD(T) interaction energies, the sources of error are:

- Single-particle basis representation of the CCSD(T) wavefunction.
- Local approximations of long-range electron correlation according to the LNO scheme.
- Neglected core electron correlation.
- Missing high-order many-electron contributions beyond CCSD(T).

- plane-wave basis (VASP)  
- canonical approach

Theory	Interaction energy
MP2	$-38.5 \pm 0.5$
CCSD	$-13.4 \pm 0.5$
CCSD(T)	$-21.1 \pm 0.5$
LNO-CCSD(T)	$-20.6 \pm 0.6$
DLPNO-CCSD(T) <sub>0</sub>	$-20.9 \pm 0.4$
PNO-LCCSD(T)-F12	-20.0
DMC	-18.1(8)
DMC	-17.5(14)

- Gaussian type basis sets  
- local approaches

# Accurate ab-initio calculations:

Interactions between large molecules pose a puzzle for reference quantum mechanical methods

Yasmine S. Al-Hamdani<sup>1,2,9</sup>, Péter R. Nagy<sup>3,9</sup>, Andrea Zen<sup>4,5,6,7</sup>, Dennis Barton<sup>2</sup>, Mihály Kállay<sup>3</sup>, Jan Gerit Brandenburg<sup>8,10</sup> & Alexandre Tkatchenko<sup>10</sup>

NATURE COMMUNICATIONS | (2021)12:3927 | <https://doi.org/10.1038/s41467-021-24119-3> | [www.nature.com/naturecommunications](http://www.nature.com/naturecommunications)

In obtaining CCSD(T) interaction energies, the sources of error are:

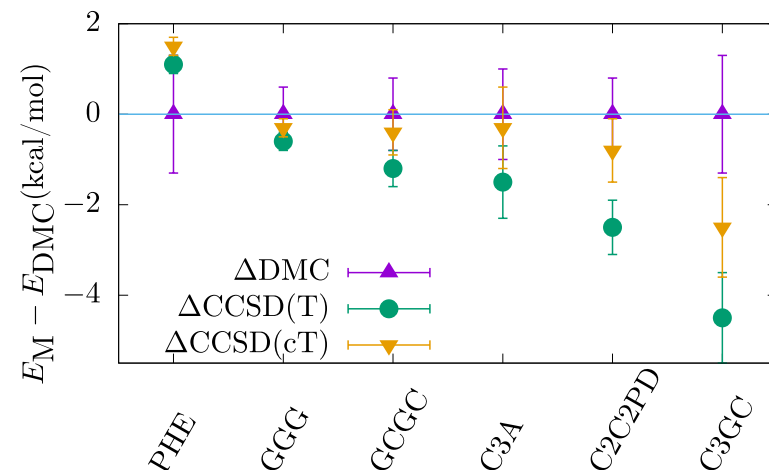
- Single-particle basis representation of the CCSD(T) wavefunction.
- Local approximations of long-range electron correlation according to the LNO scheme.
- Neglected core electron correlation.
- Missing high-order many-electron contributions beyond CCSD(T).

CCSD(cT)

## Understanding Discrepancies of Wavefunction Theories for Large Molecules

Tobias Schäfer, Andreas Irmeler, Alejandro Gallo, and Andreas Grüneis

arXiv: 2407.01442, Nat. Comm (accepted)

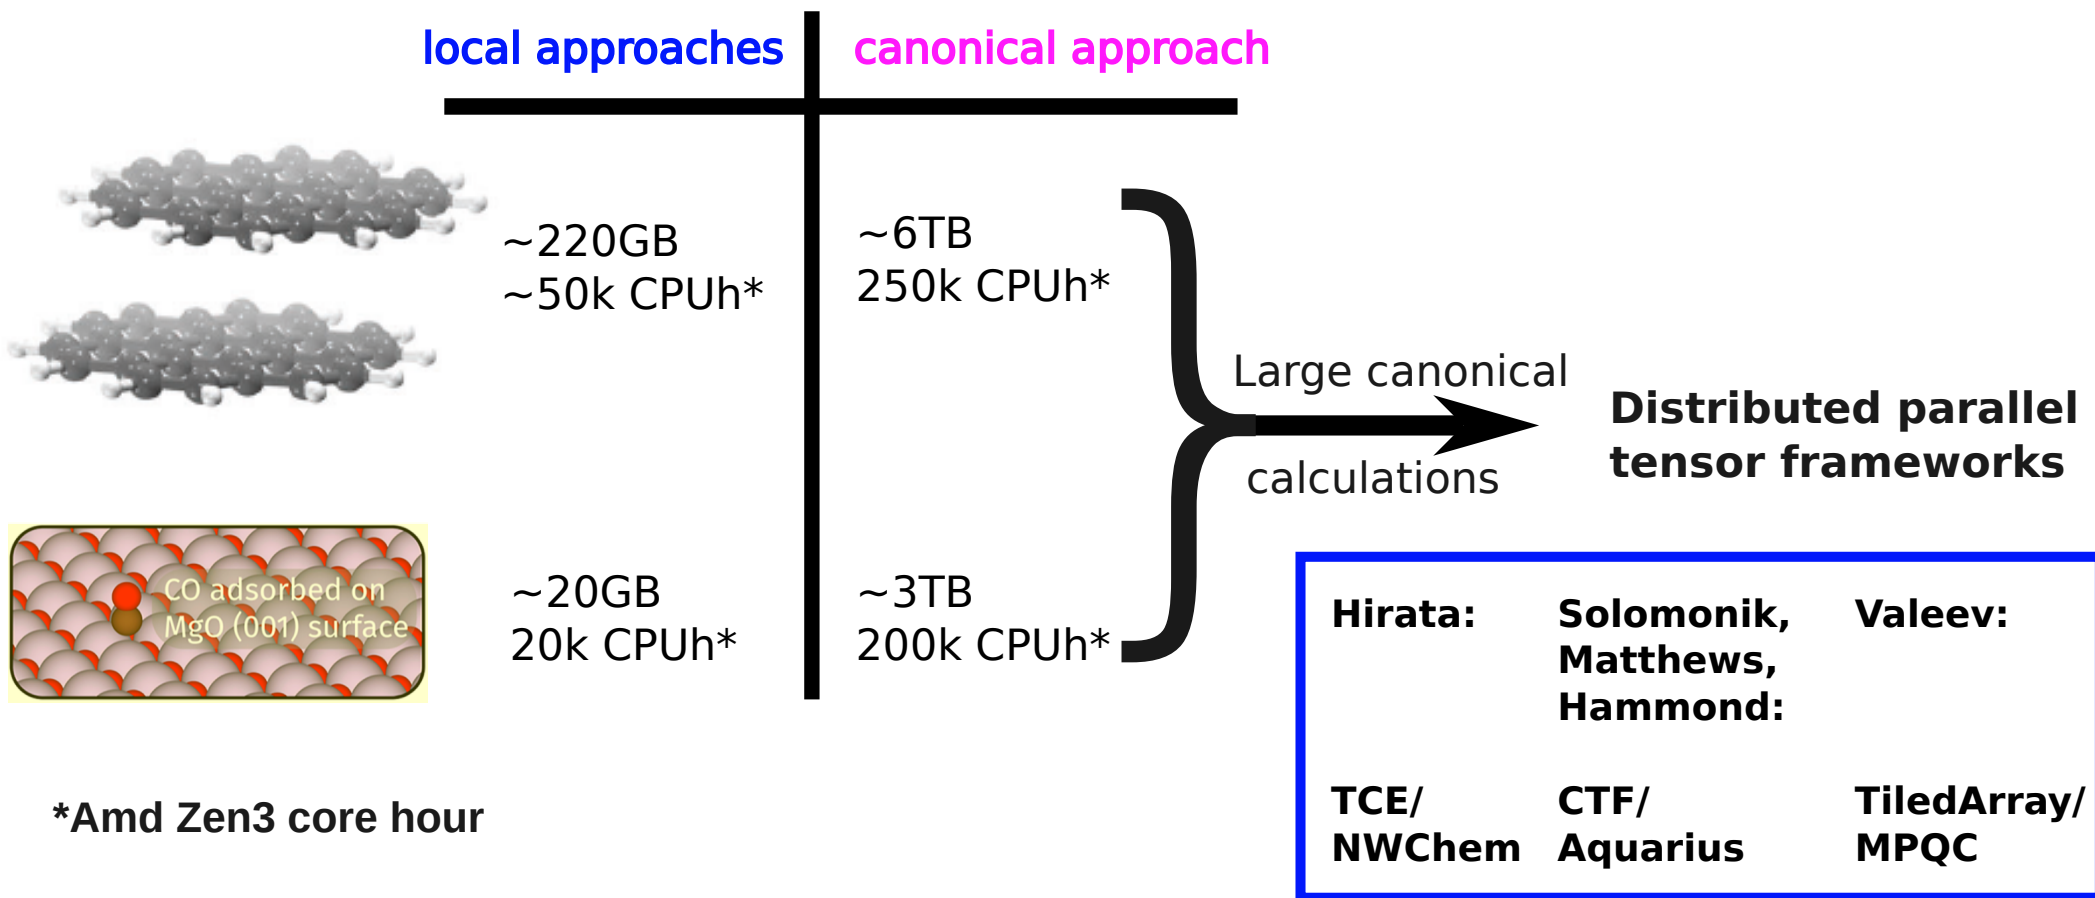


PHYSICAL REVIEW LETTERS **131**, 186401 (2023)

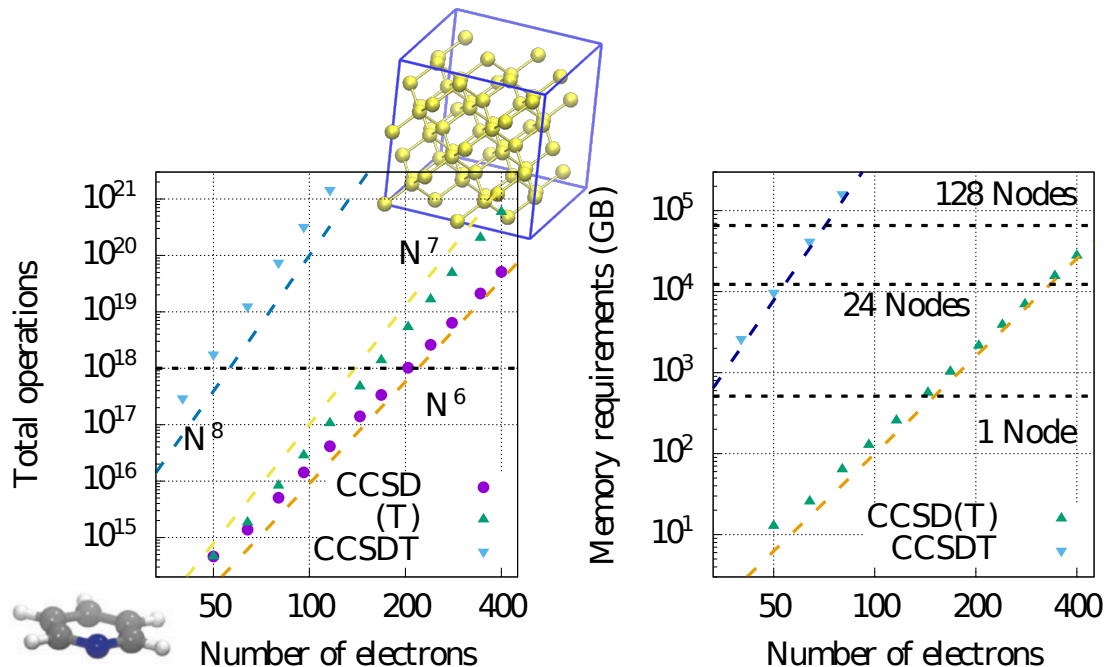
## Averting the Infrared Catastrophe in the Gold Standard of Quantum Chemistry

Nikolaos Masios, Andreas Irmeler, Tobias Schäfer, and Andreas Grüneis

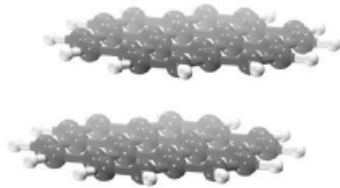
# Large-scale CCSD(T) calculations:



# Coupled-cluster methods: Computational cost and performance



$$r_{ij}^{ab} = \sum_{ck} v_{ic}^{ak} t_{kj}^{cb}$$



$$i=j=108$$

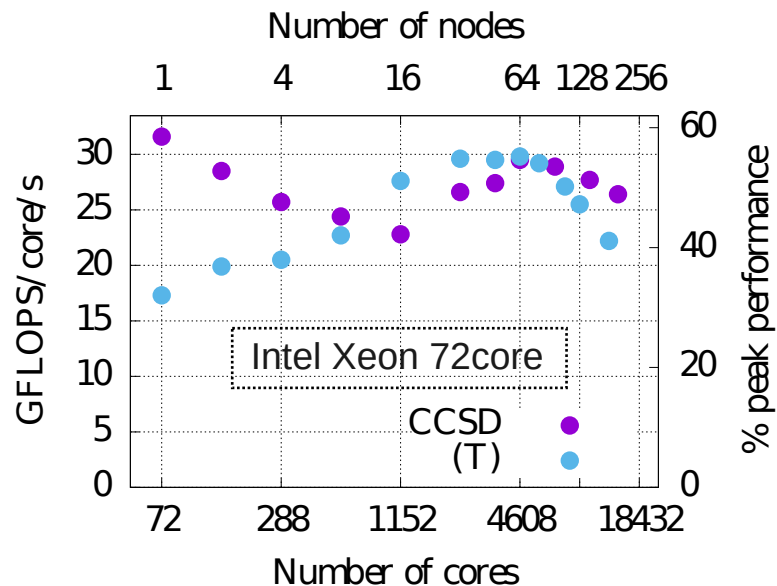
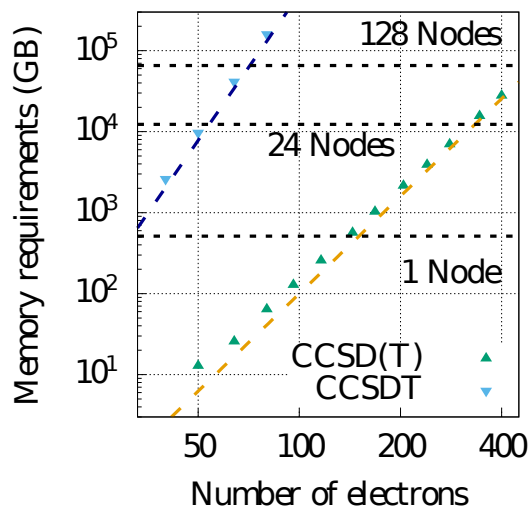
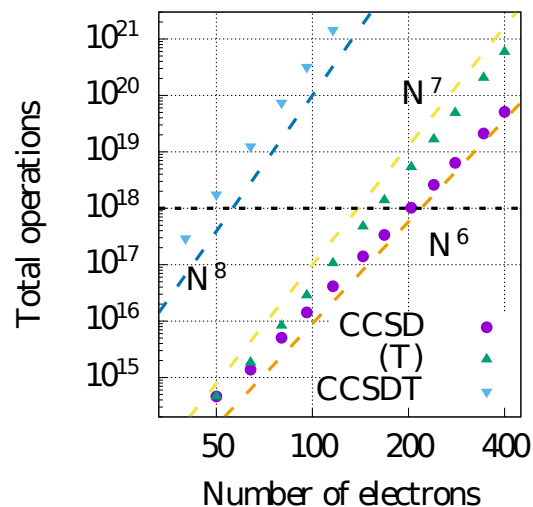
$$a=b=1164 *$$



360GB mem.  
~1h Zen3 128core

\*aug-cc-pvdz

# Coupled-cluster methods: Computational cost and performance





# cc4s: massively parallel coupled-cluster code for solids and molecules



**MRCC**



**calculate Hartree-Fock**

**write to disk**

`EigenEnergies.elements`  
`EigenEnergies.yaml`

`CoulombVertex.elements`  
`CoulombVertex.yaml`

**read from disk**

**cc4s/ctf**

**Ccsd(T)**

**Ccsd(cT)**

**Ccsdt**

**Eom-CCSD**

...

# ctf - cyclops tensor framework

- tensors are distributed over all ranks (cyclic distribution)
- transpose tensor indices (local op.)
- parallel contraction using variants of the summa algorithm

$$r_{ij}^{ab} = \sum_{ck} v_{ic}^{ak} t_{kj}^{cb}$$

[github.com/cyclops-community/ctf](https://github.com/cyclops-community/ctf)

```
#include <ctf.hpp>

int main(int argc, char ** argv){
    MPI_Init(&argc, &argv);
    int No(10), Nv(100);

    CTF::Tensor<> r(4, {Nv,Nv,No,No});
    CTF::Tensor<> v(4, {Nv,No,No,Nv});
    CTF::Tensor<> t(4, {Nv,Nv,No,No});

    r["abij"] = v["akic"] * t["bckj"];

    MPI_Finalize();
    return 0;
}
```

# ctf - cyclops tensor framework

## Simple Example

$$C_{ijkl} = \sum_G A_{ik}^G B_{jl}^G$$

Contraction performance model: find the best cyclic distribution which allows contraction:

- **redistribute** all tensors ( $N^4 + N^3$ )
- **transpose** rhs-tensors ( $N^3$ )
- parallel **contraction** via summa ( $N^5$ )
- **transpose** result tensor ( $N^4$ )
- bring result tensor in original **distribution** ( $N^4$ )

# ctf - cyclops tensor framework

## Simple Example

$$C_{ijkl} = \sum_G A_{ik}^G B_{jl}^G$$

Contraction performance model: find the best cyclic distribution which allows contraction:

- **redistribute** all tensors ( $N^4 + N^3$ )
- **transpose** rhs-tensors ( $N^3$ )
- parallel **contraction** via summa ( $N^5$ )
- **transpose** result tensor ( $N^4$ )
- bring result tensor in original **distribution** ( $N^4$ )

Kernel	Relative time
<b>redistribute</b>	5-30%
<b>transpose</b>	0-5%
<b>gemm</b>	40-80%
<b>bcast/reduce</b>	15-40%

# ctf - cyclops tensor framework

## Simple Example

$$C_{ijkl} = \sum_G A_{ik}^G B_{jl}^G$$

Contraction performance model: find the best cyclic distribution which allows contraction:

- **redistribute** all tensors ( $N^4 + N^3$ )
- **transpose** rhs-tensors ( $N^3$ )
- parallel **contraction** via summa ( $N^5$ )
- **transpose** result tensor ( $N^4$ )
- bring result tensor in original **distribution** ( $N^4$ )

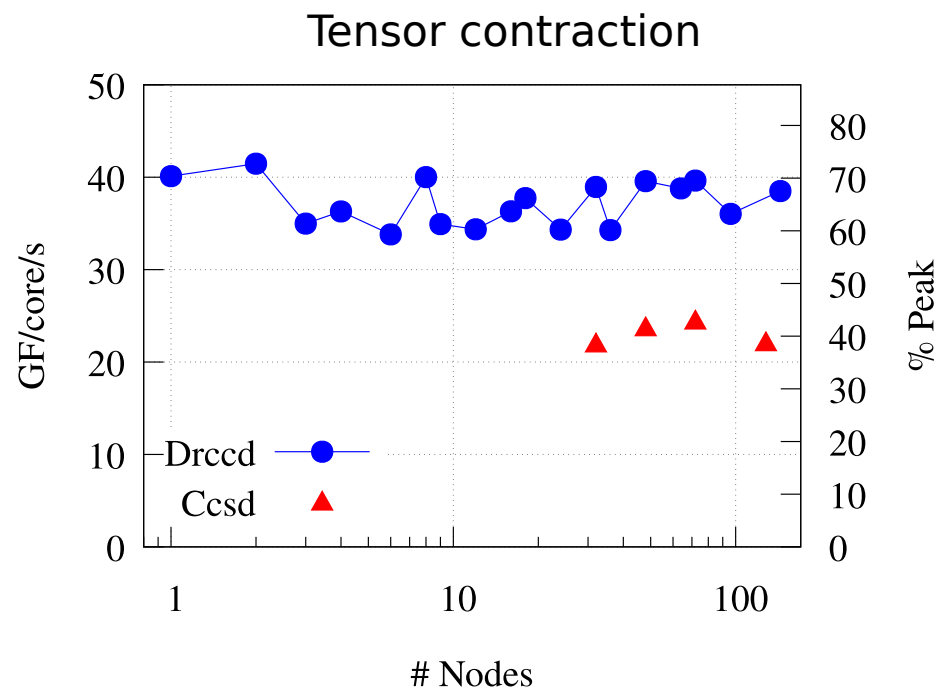
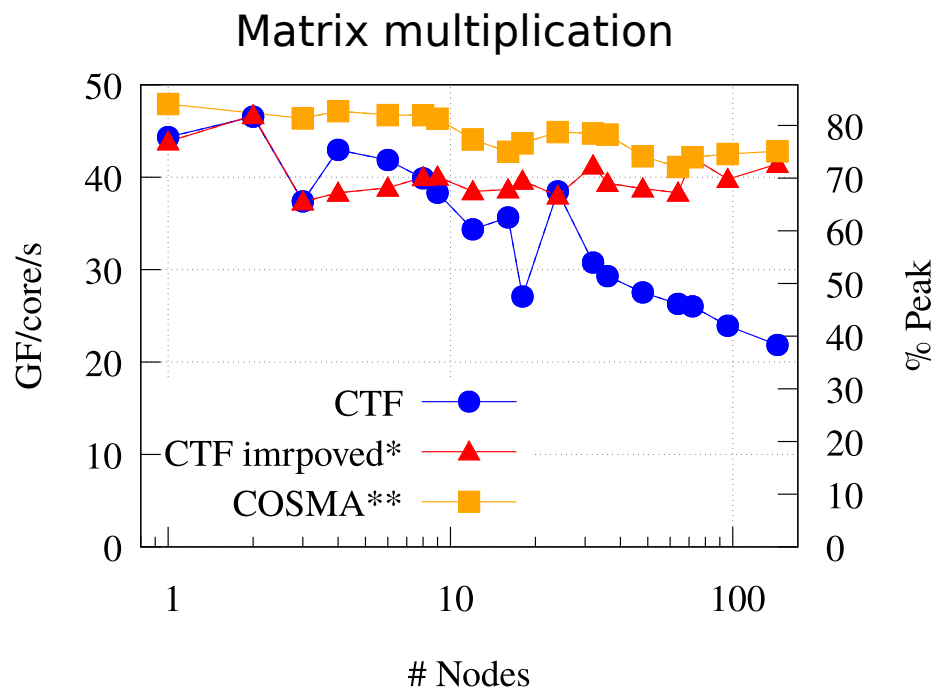
Kernel	Relative time
<b>redistribute</b>	5-30%
<b>transpose</b>	0-5%
<b>gemm</b>	40-80%
<b>bcast/reduce</b>	15-40%



GEMM dominates runtime and runs at peak efficiency

# ctf - cyclops tensor framework

## Performance



\* Irmeler, ..., Solomonik; EuroPar 23'

\*\*Kwasniewski, ..., Hoefler; SC'19

# ctf: cyclic distribution

0	3	6	0	3	6
1	4	7	1	4	7
2	5	8	2	5	8
0	3	6	0	3	6
1	4	7	1	4	7
2	5	8	2	5	8

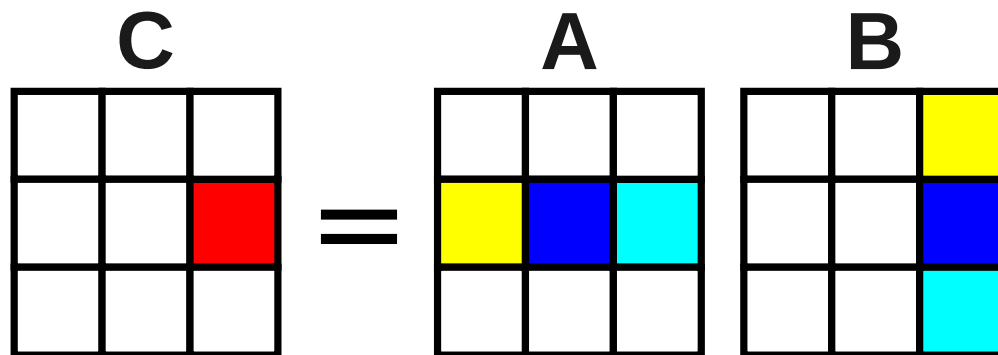
cyclic phase among all  
edges of each tensor

# ctf: cyclic distribution

0	3	6	0	3	6
1	4	7	1	4	7
2	5	8	2	5	8
0	3	6	0	3	6
1	4	7	1	4	7
2	5	8	2	5	8

cyclic phase among all  
edges of each tensor

SUMMA algorithm:





# ctf: cyclic distribution

0	3	6	0	3	6
1	4	7	1	4	7
2	5	8	2	5	8
0	3	6	0	3	6
1	4	7	1	4	7
2	5	8	2	5	8

cyclic phase among all  
edges of each tensor

In a tensor contraction - the phases  
of the tensors in the contraction  
have to be consistent:


$$\begin{bmatrix} \text{orange} & \text{red} & \text{green} \\ \text{blue} & \text{yellow} & \text{purple} \end{bmatrix} = \begin{bmatrix} \text{orange} & \text{red} & \text{green} \\ \text{blue} & \text{yellow} & \text{purple} \end{bmatrix} \times \begin{bmatrix} \text{orange} & \text{red} & \text{green} \\ \text{blue} & \text{yellow} & \text{purple} \end{bmatrix}$$

 **Redistribute!**



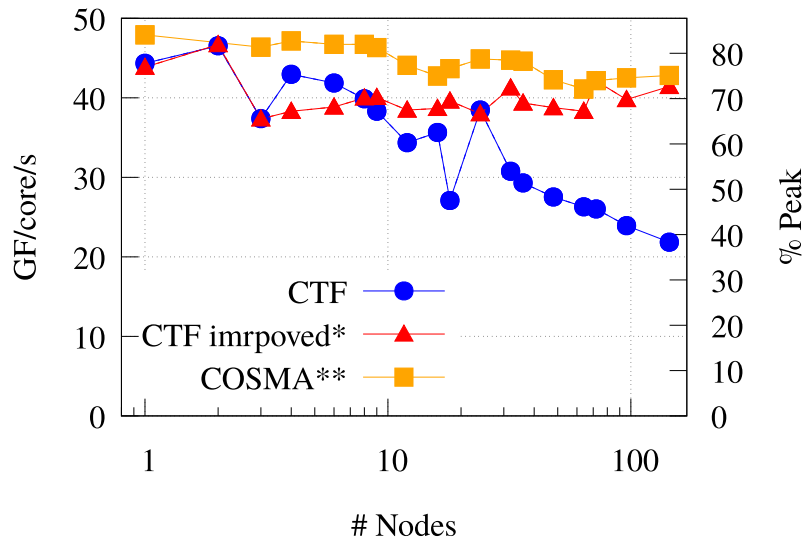
In practice, redistribution is  
required in every contraction!!  
(true for CC calculations)!

# ctf - cyclops tensor framework

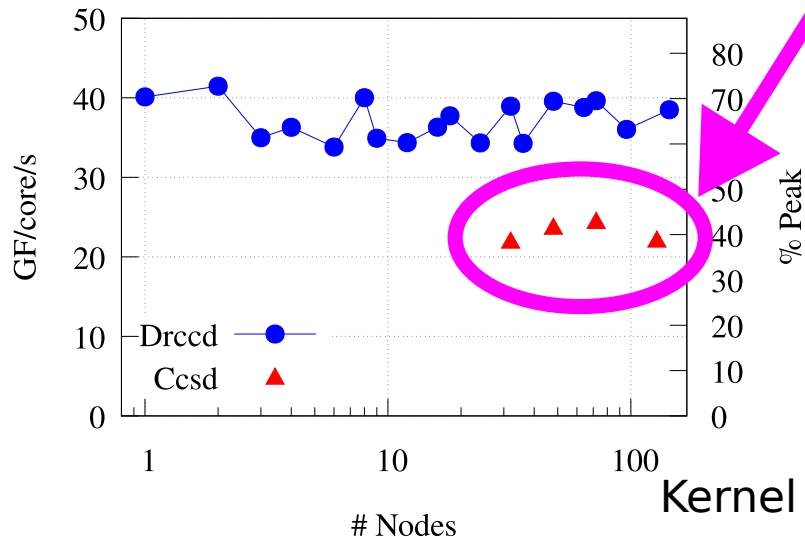
## Performance

1000 M gemm operations  
VS.  
M redistribute/transpose

Matrix multiplication



Tensor contraction



$N^3$  gemm operations  
VS.  
 $N^2$  transpose/redistribute

Kernel	Relative time
redistribute	5-30%
transpose	0-5%
gemm	40-80%
bcast/reduce	15-40%

# ctf: symmetries

$$t_{ij}^{ab} = -t_{ji}^{ab} = -t_{ij}^{ba} = t_{ji}^{ba}$$

```
int asas[] = {AS, NS, AS, NS};
int abij[] = {nocc, nocc, nvir, nvir};
Tensor<> T(4, abij, asas);
Tensor<> R(4, abij, asas);
{
    int abcd[] = {nvir, nvir, nvir, nvir};
    Tensor<> V(4, abcd, asas);
    R["abij"] = V["abcd"] * T["cdij"];
}
```



Often, we cannot store  $V["abcd"]$  in memory

for  $x$  in blocks( $a$ ),  $y$  in blocks( $b$ ) :

$$V["xycd"] = X["Gxc"] * X["Gyd"];$$

$$R["xyij"] += V["xycd"] * T["cdij"];$$

# ctf: symmetries

$$t_{ij}^{ab} = -t_{ji}^{ab} = -t_{ij}^{ba} = t_{ji}^{ba}$$

```
int asas[] = {AS, NS, AS, NS};
int abij[] = {nocc, nocc, nvir, nvir};
Tensor<> T(4, abij, asas);
Tensor<> R(4, abij, asas);
{
    int abcd[] = {nvir, nvir, nvir, nvir};
    Tensor<> V(4, abcd, asas);
    R["abij"] = V["abcd"] * T["cdij"];
}

{
    int nsns[] = {NS, NS, NS, NS};
    int akic[] = {nvir, nocc, nocc, nvir};
    Tensor<> V(4, akic, nsns);
    R["abij"] = V["akic"] * T["cbkj"];
}
```

Often, we cannot store  $V["abcd"]$  in memory



for  $x$  in blocks( $a$ ),  $y$  in blocks( $b$ ) :

$$V["xycd"] = X["Gxc"] * X["Gyd"];$$

$$R["xyij"] += V["xycd"] * T["cdij"];$$



- transform  $T$  to non-symmetric tensor
- contraction without symmetry
- transform  $R$  to symmetric tensor

# ctf: porting to GPUs

Contraction performance model: find the best cyclic distribution which allows contraction:

- **redistribute** all tensors ( $N^4 + N^3$ )
  - **transpose** rhs-tensors ( $N^3$ )
  - parallel **contraction** via summa ( $N^5$ )
  - **transpose** result tensor ( $N^4$ )
  - bring result tensor in original **distribution** ( $N^4$ )
- } Relatively straightforward to port to GPUs.



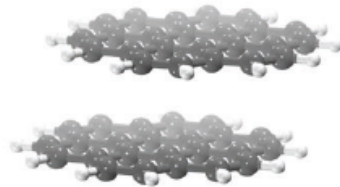
In practice, redistribution is required in every contraction!! (true for CC calculations)!



Redistribution step can not be simply overlapped with other parts of the calculation.

# Atrip: evaluation of (T)

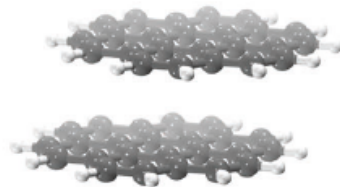
$$\delta E_{(T)} = \sum_{ijk} \sum_{abc} \hat{t}_{ijk}^{abc} \bar{t}_{ijk}^{abc}$$



$$t_{ijk}^{abc} \approx 10000\text{TB}$$

# Atrip: evaluation of (T)

$$\delta E_{(T)} = \sum_{ijk} \sum_{abc} \hat{t}_{ijk}^{abc} \bar{t}_{ijk}^{abc}$$



$$t_{ijk}^{abc} \approx 10000\text{TB}$$

Instead:

Solve CCSD to get  $T_1, T_2$

for occupied  $(i, j, k)$  :

$\hat{t}^{abc} = \text{build\_intermediate}(T_1, T_2, \text{integrals})$  #for fixed  $ijk$

$\bar{t}^{abc} = \text{build\_intermediate}(T_1, T_2, \text{integrals})$

$$\delta E_{(T)} \leftarrow \sum_{abc} \hat{t}^{abc} \bar{t}^{abc}$$



Not efficiently  
implementable in CTF



Memory footprint does not exceed  
that of CCSD

# Conclusions:

- CTF allows efficient CCSD calculations up to 100 nodes
- Careful assessment of whether symmetries are efficient
- Not all coupled-cluster variants can be efficiently implemented with CTF
- Efficiently porting CTF to GPUs is challenging